


Homework 01

Problem 1:

Homework 01: Answer



Known: $P(w_2) = 4/9$, $P(w_1) = 5/9$
 Since we have 2 separate classes
 $P(w_2 \cup w_1) = 1$

from def. of uniform distribution
 $P(x|w_2) = \frac{1}{6-2} = 1/4$

• P_e for class w_2 :

$$P_{e2} = \int_{-\infty}^T P(x|w_2)P(w_2) = \int_2^T \left(\frac{1}{4}\right)\left(\frac{4}{9}\right) dx = \frac{x}{9} \Big|_{x=2}^{x=T} = \frac{(T-2)}{9}$$

NOTE: since $P_{e2} \geq 0$, T lower bounded by 2

• P_e for class w_1 :

$$P_{e1} = \int_T^{\infty} P(x|w_1)P(w_1) = \int_T^4 \left(-\frac{1}{4}x + 1\right)\left(\frac{5}{9}\right) dx$$

$$= \left(-\frac{5x^2}{72} + \frac{5x}{9}\right) \Big|_{x=T}^{x=4} = \frac{5T^2 - 80}{72} + \frac{20 - 5T}{9}$$

NOTE: threshold T bounded by the max value w_1 can take

• Finding P_e :

• P_e is total probability error of classes $\{w_1, w_2\}$

P_e for both classes ≥ 0

$P_{e2} \geq 0 \rightarrow T \geq 2$

$P_{e1} \geq 0 \rightarrow T \leq 4$

$$P_e = P_{e1} + P_{e2} = \frac{5T^2 - 32T + 64}{72}$$

$$\frac{\partial P_e(T)}{\partial T} = 10T - 32 = 0$$

$$= T = 32/10$$

taking derivative
 w.r.t T to find
 T which minimizes
 P_e

Resulting P_e using T which
 minimizes P_e :

$$T = 32/10 \quad P_e = 0.178$$

when $T = 4$, $P_e = 0.222$
 when $T = 2$, $P_e = 0.278$

Figure 1: Handwritten answer of the problem 1

Problem 2:

Part 2.1:

```
1  # Calculate mean of the features
2  mean_face = np.mean(x_train, axis=0)
3  mean_face = np.reshape(mean_face, image_shape)
4
5  # Obtain eigenfaces via Linear dimension reduction
6  # Use Singular Value Decomposition of the data to project it to a lower dim space.
7  N = x_train.shape[0]
8  pca = PCA(n_components = N-1)
9  # Normalize training images via subtracting the mean image
10 pca_fit = pca.fit(np.subtract(x_train, mean_face.reshape(-1)))
11
12 # Fit the model onto the data
13 eigenface1 = pca_fit.components_[0].reshape(image_shape)
14 eigenface2 = pca_fit.components_[1].reshape(image_shape)
15 eigenface3 = pca_fit.components_[2].reshape(image_shape)
16 eigenface4 = pca_fit.components_[3].reshape(image_shape)
```




Figure 2: Plot of the mean face and first 4 eigenfaces

Part 2.2 and Part 2.3:

```
1  # Normalize the image by subtracting the mean face from original face
2  # Project the data onto the PCA space.
3  pca_orig_image = pca_fit.transform(np.subtract(original_face_flat,
4  mean_face.reshape(-1)))
5
6  # Use PCA space dimensions 3, 45, 140 and 229 to reconstruct images
7  for i in (3, 45, 140, 229):
8      img_pca = np.dot(pca_orig_image[0,:i], pca_fit.components_[i])
9      + mean_face.reshape(-1)
10     # Mean squared error between original and the reconstructed image
11     mse = mean_squared_error(img_pca.reshape((1, img_pca.shape[0])),
12     original_face_flat)
13     img_pca = img_pca.reshape(original_face.shape)
14     plt.subplot(idx)
```

```

15 plt.title("n=%s, mse=%.5f" % (i, mse))
16 plt.imshow(img_pca, cmap='gray')
17 idx += 1

```




Figure 3: Reconstructed images using dimensions N equals to 3, 45, 140 and 229

Part 2.4:

```

1  # Normalize and project the training data onto the PCA subspace
2  x_train_normalized = pca_fit.transform(np.subtract(x_train, mean_face.reshape(-1)))
3  y_train = np.array(y_train)
4
5  # Configure k-Nearest Neighbours Model with k equals to 1, 3 and 5
6  params = {'n_neighbors': [1, 3, 5]}
7  kNN = KNeighborsClassifier()
8  # Perform 3 fold cross validation
9  cv = GridSearchCV(kNN, params, cv = 3)
10
11 df = dict()
12 for i in (3, 45, 140):
13     cv.fit(x_train_normalized[:, :i], y_train)
14     # Show cross validation results in terms of mean test score
15     df['n = ' + str(i)] = np.array(cv.cv_results_['mean_test_score'])

```

N	k = 1	k = 3	k = 5
3	0.691667	0.595833	0.525000
45	0.941667	0.841667	0.750000
140	0.954167	0.829167	0.745833

Part 2.5:

I decided to choose $N = 45$ and $k = 1$, even though it performs worse than $N = 140$ and $k = 1$. The reason behind my logic is that even though we increase dimension by 211.1% percent we only get 1.327% increase in the test accuracy which is not significant and might be the indication of overfitting to the data. I obtained **Accuracy on the test set: 0.95625**.

```

1  # best parameters
2  k = 1
3  n = 45
4
5  # Project training images onto the PCA subspace
6  pca_test = pca_fit.transform(np.subtract(x_test, mean_face.reshape(-1)))
7
8  x_train_normalized = pca_fit.transform(np.subtract(x_train, mean_face.reshape(-1)))
9  y_train = np.array(y_train)
10
11 # train and evaluate the k-Neighbours model using the chosen parameters
12 kNN_opt = KNeighborsClassifier(n_neighbors = k)
13 kNN_opt.fit(x_train_normalized[:, :n], y_train)
14 pred = kNN_opt.predict(pca_test[:, :n])
15 print("Accuracy on the test set:", accuracy_score(y_pred = pred, y_true = y_test))

```

Question 3:

Part 3.1:

I couldn't describe an image by just seeing random patches from an image. Only cases I got correct are the ones where I can identify definite shapes such as yellow edges of the bananas. Also knowing that I only have 4 classes make my guess more easier since I have 25% correct guess possibility with a random guess. On top of that, I did a quick look over the dataset which make my guesses easier but since I am overfitting/memorizing the dataset it is not an indication of being able to identify images by just looking at image patches.




Figure 4: 3 random banana patches




Figure 5: 3 random fountain patches




Figure 6: 3 random reef patches




Figure 7: 3 random tractor patches

Part 3.2:

```

1  # Converting image patches to 1-d and storing
2  X_train_patches = []
3  X_test_patches = []
4
5  # Note: X_train_patch contains non flattened image patches
6  for img in X_train_patch:
7      for patch in img:

```

```

8         patch = patch.reshape(-1)
9         X_train_patches.append(patch)
10
11     for img in X_test_patch:
12         for patch in img:
13             patch = patch.reshape(-1)
14             X_test_patches.append(patch)
15
16     X_train_patches = np.array(X_train_patches)
17     X_test_patches = np.array(X_test_patches)
18
19     print("X_train_patches shape: ", X_train_patches.shape)
20     # X_train_patches shape: (24000, 768)
21     print("X_test_patches shape: ", X_test_patches.shape)
22     # X_test_patches shape: (8000, 768)
23
24     # Configure k-means with C = 15 and maximum iteration = 5000
25     kmeans = KMeans(n_clusters=15, random_state=0, max_iter=5000).fit(X_train_patches)
26
27     centroids = kmeans.cluster_centers_
28     kmeans_clusters = kmeans.labels_
29
30     # Construction of 3 dimensional PCA subspace
31     pca = PCA(n_components=3)
32     pca_X_train_patches = pca.fit_transform(X_train_patches)
33
34     # get 6 random clusters
35     sample_C = 6
36     cluster_samples = np.random.choice(np.arange(15), size=sample_C, replace=False)
37
38     # get the centroids of the sampled clusters
39     centroid_samples = centroids[cluster_samples]
40     pca_centroids = pca.transform(centroid_samples)

```

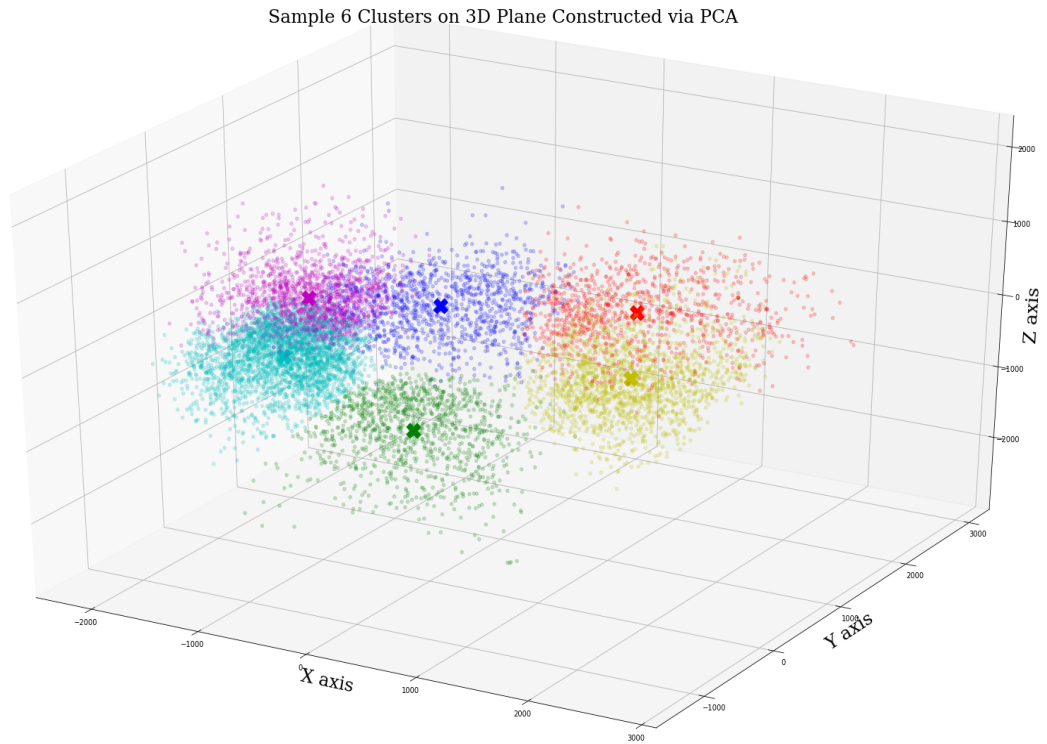


Figure 8: Visual words together with their associated features on the PCA subspace

Part 3.3:

Visualizing histogram of Softmax images for each category of images. The histograms created by constructing a Softmax matrix by calculating the euclidean distance between each cluster centroids and the training images. After reciprocal of the euclidean distance normalized, BoW calculated via taking maximum of the each encoded features in that dimension.

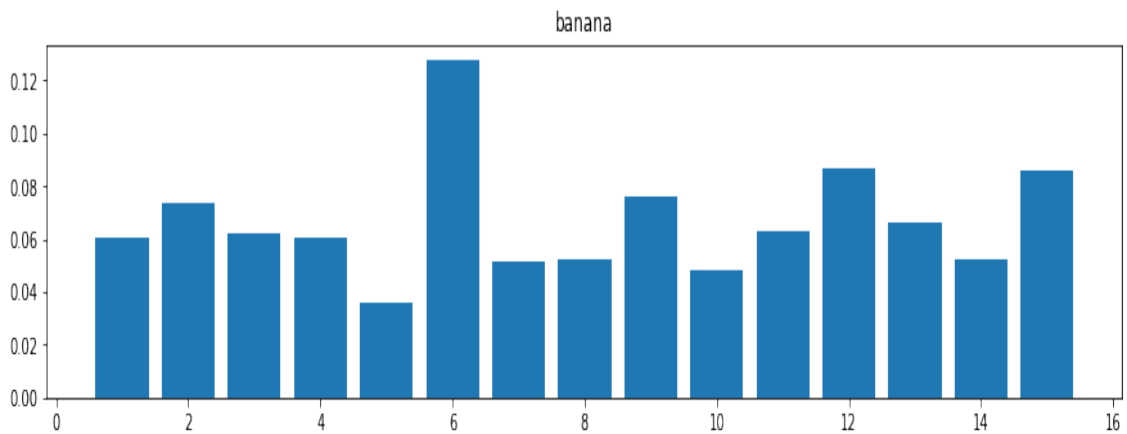


Figure 9: Histogram of the BoW: Banana

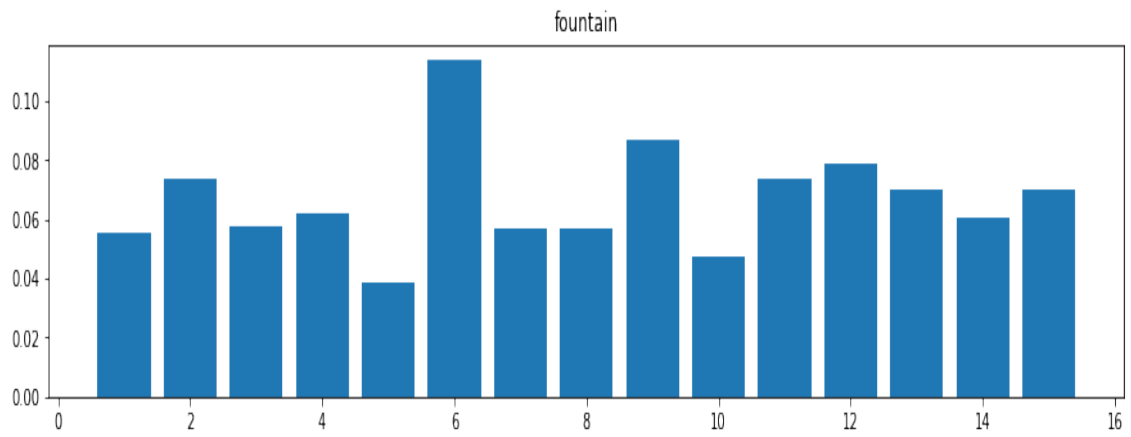


Figure 10: Histogram of the BoW: Fountain

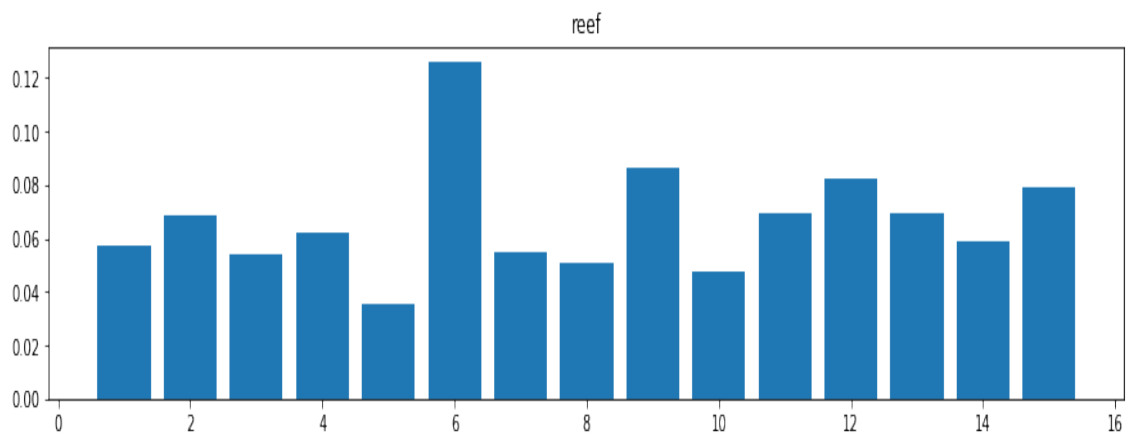


Figure 11: Histogram of the BoW: Reef

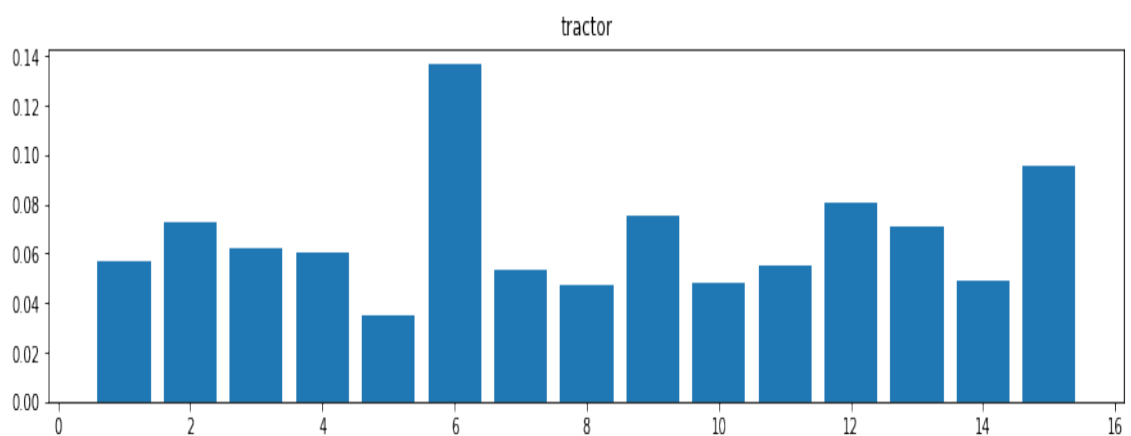


Figure 12: Histogram of the BoW: Tractor

Part 3.4:
Accuracy of Softmax on the test set: 0.552


```

1  # Configure kNN with Number of neighbours is 5
2  knn = KNeighborsClassifier(n_neighbors=5)
3  knn.fit(X_train_softmax, Y_train)
4
5  print(X_test.shape)
6  # (500, 12288)
7
8  # Compute softmax of the test set
9  X_test_softmax = Softmax(X_test, centroids)
10
11 print("Accuracy of Softmax on test set: ",
12       accuracy_score(y_pred = knn.predict(X_test_softmax), y_true=Y_test))

```

Collaborators

I used Andrew Ng's Coursera Machine Learning Course as a reference and Scikit-learn documentation. Other than that no collaborators.

1. <https://scikit-learn.org/stable/documentation.html>